

RAG-Guard Security & Document Ingestion Test

This PDF document is generated to test PyMuPDF (fitz) text parsing capabilities.

Security Concepts:

1. Prompt Injection: Traditional prompt injection occurs when a user directly crafts instructions to override the system prompt. In contrast, 'Indirect Prompt Injection' is a new vulnerability where the attack payload is stored inside an external data source (like a PDF or webpage). When the RAG pipeline retrieves this document, the payload enters the context window and tricks the LLM into executing malicious commands, ignoring its original system instructions.
2. PII Detection: High-risk documents often contain Personally Identifiable Information (PII) like names, addresses, phone numbers, or social security numbers. Presidio scanner intercepts these chunks pre-generation.
3. Local Model Verification: Milestone 1 is configured to use gemma3:1b via local Ollama.